


Compute / communication overlap on the GPU timeline (rows are CUDA-stream groups; bars are kernels)

rank 0 (PP stage 0) — one full step (5373 ms). comm bar with no 'compute' bar at the same time = exposed



rank 16 (PP stage 1) — one full step (5373 ms). comm bar with no 'compute' bar at the same time = exposed

